Data Processing in Python

The Modern Python Data Stack

Somto Muotoe

January 2026

Agenda

Part 1: File Operations

- What is pathlib?
- Path basics & properties
- Reading & writing files
- Directory operations

Part 2: Pandas

- DataFrame fundamentals
- Reading data files
- Data manipulation
- Aggregations & grouping

Part 3: Polars

- Why Polars over Pandas?
- Lazy vs Eager evaluation
- Expression API
- Performance comparison

Part 4: DuckDB

- What is DuckDB?
- SQL on DataFrames
- When to use DuckDB
- Integration patterns

Setup: Clone the repo for data files and exercises → github.com/smuotoe/mcda-data-processing-tutorial

Part 1

File Operations with pathlib

What is pathlib?

A **built-in Python module** for working with file system paths in an object-oriented way. Part of the standard library since Python 3.4.

The Problem

- File paths differ across operating systems (`/` on Unix vs `\` on Windows)
- String manipulation for paths is error-prone
- `os.path` functions are verbose and scattered
- Edge cases are easy to miss (e.g., `data/` vs `data`)

How pathlib Solves It

- **Cross-platform** - handles path separators automatically
- **Object-oriented** - paths are objects with methods
- **Intuitive operators** - use `/` to join paths
- **All-in-one** - file operations built into Path objects

```python
from pathlib import Path

path = Path("data") / "sales.csv"
content = path.read_text()
```

Why pathlib?

The Old Way: os.path

```python
import os

# Joining paths - error prone
path = os.path.join("data", "sales", "2024.csv")

# Reading a file
with open(path, 'r') as f:
    content = f.read()

# Getting file info
exists = os.path.exists(path)
is_file = os.path.isfile(path)
name = os.path.basename(path)
```

String-based, verbose, platform issues

The Modern Way: pathlib

```python
from pathlib import Path

# Joining paths - clean and safe
path = Path("data") / "sales" / "2024.csv"

# Reading a file
content = path.read_text()

# Getting file info
exists = path.exists()
is_file = path.is_file()
name = path.name
```

Object-oriented, readable, cross-platform

An edge case pathlib handles: `"data/" + "file.txt"` → `"data/file.txt"` but `"data" + "file.txt"` → `"datafile.txt"` ✗

With pathlib: `Path("data") / "file.txt"` and `Path("data/") / "file.txt"` both → `data/file.txt` ✓

Path Basics

Creating Paths

```python
from pathlib import Path

# Current directory
cwd = Path.cwd()

# Home directory
home = Path.home()  # C:\Users\somto

# From string
data_path = Path("data/sales")

# Absolute path
abs_path = Path("C:\Users\somto\projects")

# Joining paths with /
file_path = data_path / "report.csv"
```

Path Properties

```python
path = Path("data/sales/2024_q1.csv")

path.name          # "2024_q1.csv"
path.stem          # "2024_q1"
path.suffix        # ".csv"
path.parent        # Path("data/sales")
path.parts         # ("data", "sales", "2024_q1.csv")

# Resolve to absolute
path.resolve()   # Full absolute path

# Check path type
path.is_file()   # True/False
path.is_dir()    # True/False
path.exists()    # True/False
```

Reading & Writing Files

Reading Files

```python
from pathlib import Path

path = Path("data/config.json")

# Read entire file as text
content = path.read_text()

# Read as bytes (for binary files)
binary = path.read_bytes()

# With encoding
content = path.read_text(encoding="utf-8")
```

Writing Files

```python
from pathlib import Path
import json

path = Path("output/results.json")

# Ensure parent directory exists
path.parent.mkdir(parents=True, exist_ok=True)

# Write text
path.write_text("Hello, World!")

# Write JSON data
data = {"name": "test", "value": 42}
path.write_text(json.dumps(data, indent=2))

# Write bytes
path.write_bytes(b"binary data")
```

Tip: `read_text()` and `write_text()` handle file opening/closing via context managers automatically

Directory Operations

```python
from pathlib import Path

# Create directories
Path("output/reports").mkdir(parents=True, exist_ok=True)

# List directory contents
for item in Path("data").iterdir():
    print(item.name, "- dir" if item.is_dir() else "- file")

# Find files with glob patterns
csv_files = list(Path("data").glob("*.csv"))          # Direct children only
all_csvs = list(Path("data").rglob("*.csv"))          # All subdirectories

# Pattern matching examples
Path(".").glob("*.py")              # All Python files
Path(".").glob("**/*.json")         # All JSON files (recursive)
Path(".").glob("data_202[0-9].csv") # data_2020.csv through data_2029.csv
```

Note: `rglob("*.csv")` is shorthand for `glob("**/*.csv")` - both search recursively

Practical Example: Organizing Files

```python
from pathlib import Path
import shutil

def organize_downloads():
    downloads = Path.home() / "Downloads"
    categories = {
        "images": [".jpg", ".png", ".gif"],
        "documents": [".pdf", ".docx", ".txt"],
        "data": [".csv", ".json", ".parquet"],
    }
    for path in downloads.iterdir():  # iterate over directory contents
        if not path.is_file():
            continue
        for category, extensions in categories.items():
            if path.suffix.lower() in extensions:
                dest_dir = downloads / category
                dest_dir.mkdir(exist_ok=True)
                shutil.move(path, dest_dir / path.name)  # move to category folder
                break
```

Combines `iterdir()`, `is_file()`, `suffix`, `mkdir()`, and path joining

Exercise: File Operations

Task: Create a Student Profile System

1. Create a directory named `profiles`

2. Save your profile as JSON (name, languages, hobby)

3. List all `.json` files in the directory

4. Read and display your profile

```python
from pathlib import Path
import json


def quick_profile_update():
    # 1. Create profiles directory
    # TODO: Use Path().mkdir()

    # 2. Create and save profile
    my_profile = {"name": "Your Name", "languages": ["Python"]}
    # TODO: Write to profiles/my_profile.json

    # 3. List all profiles
    # TODO: Use glob("*.json")

    # 4. Read and print
    # TODO: Use read_text() and json.loads()
```

Part 2

Pandas: Fundamentals

What is Pandas?

A Python library for **data manipulation and analysis**, providing fast, flexible data structures for working with tabular (spreadsheet-like) data.

The De Facto Standard

- Created in 2008 by Wes McKinney
- Built on NumPy
- ~95% of Python data analysis uses pandas
- Massive ecosystem and community
- Excellent documentation

Core Data Structures

- **Series**: 1D labeled array
- **DataFrame**: 2D labeled table

```python
import pandas as pd

# Series: 1D labeled array
prices = pd.Series([1.20, 0.50, 0.80])
# 0    1.20
# 1    0.50
# 2    0.80

# DataFrame: 2D labeled table
df = pd.DataFrame({
    "product": ["Apple", "Banana", "Orange"],
    "price": [1.20, 0.50, 0.80],
})
#    product  price
# 0    Apple   1.20
# 1   Banana   0.50
# 2   Orange   0.80
```

Reading Data

Common File Formats

```python
import pandas as pd
from pathlib import Path

data = Path("data")

# CSV files
df = pd.read_csv(data / "sales.csv")

# Excel files
df = pd.read_excel(data / "report.xlsx")

# JSON files
df = pd.read_json(data / "records.json")

# Parquet files (columnar format)
df = pd.read_parquet(data / "large_data.parquet")
```

Quick Data Inspection

```python
# First/last rows
df.head()        # First 5 rows
df.tail(10)      # Last 10 rows

# Shape and info
df.shape         # (rows, columns)
df.info()        # Column types, memory
df.describe()    # Statistics

# Column names and types
df.columns       # Column names
df.dtypes        # Data types

# Sample data
df.sample(5)     # Random 5 rows
```

Selecting Data

Column Selection

```python
# Single column (returns Series)
df["product"]
df.product          # Dot notation

# Multiple columns (returns DataFrame)
df[["product", "price"]]
```

Row Selection

```python
# By index position
df.iloc[0]          # First row
df.iloc[0:5]        # First 5 rows
df.iloc[[0, 2, 4]]  # Specific rows

# By label
df.loc[0]           # Row with index 0
df.loc[0:5]         # Rows 0 through 5
```

Filtering

```python
# Boolean filtering
df[df["price"] > 10]

# Multiple conditions
df[(df["price"] > 10) & (df["quantity"] > 50)]

# Using query (cleaner syntax)
df.query("price > 10 and quantity > 50")

# Check if value in list
df[df["product"].isin(["Apple", "Banana"])]
```

Combined Selection

```python
# Rows and columns together
df.loc[0:5, ["product", "price"]]
df.iloc[0:5, 0:2]
```

Data Manipulation

Creating New Columns

```python
# Direct assignment
df["revenue"] = df["quantity"] * df["price"]

# Using assign (returns new DataFrame)
df = df.assign(
    revenue=df["quantity"] * df["price"],
    tax=df["price"] * 0.1
)

# Conditional values
df["size"] = df["quantity"].apply(
    lambda x: "large" if x > 100 else "small"
)

# Using np.where
import numpy as np
df["size"] = np.where(
    df["quantity"] > 100, "large", "small"
)
```

Modifying Data

```python
# Rename columns
df = df.rename(columns={
    "old_name": "new_name"
})

# Drop columns
df = df.drop(columns=["unwanted_col"])

# Fill missing values
df["price"] = df["price"].fillna(0)

# Replace values
df["status"] = df["status"].replace({
    "Y": "Yes", "N": "No"
})

# Sort
df = df.sort_values("price", ascending=False)
```

Aggregations & Grouping

Basic Aggregations

```python
# Single column
df["price"].sum()
df["price"].mean()
df["price"].max()
df["quantity"].nunique()  # Unique count

# Multiple aggregations
df["price"].agg(["sum", "mean", "std"])

# Across DataFrame
df[["price", "quantity"]].sum()
```

GroupBy Operations

```python
# Group and aggregate
df.groupby("product")["revenue"].sum()

# Multiple aggregations
df.groupby("product").agg({
    "revenue": ["sum", "mean"],
    "quantity": "sum",
    "customer_id": "nunique"
})

# Group by multiple columns
df.groupby(["product", "region"]).sum()

# Transform (keeps original shape)
df["pct_of_total"] = (
    df["revenue"] /
    df.groupby("product")["revenue"].transform("sum")
)
```

Writing Data

```python
from pathlib import Path

output = Path("output")
output.mkdir(exist_ok=True)

df.to_csv(output / "results.csv", index=False)          # CSV - most common
df.to_excel(output / "report.xlsx", sheet_name="Sales") # Excel
df.to_parquet(output / "data.parquet")                  # Parquet - compressed, fast
df.to_json(output / "data.json", orient="records")      # JSON
```

CSV - Universal format, human-readable, opens in Excel/Google Sheets. Use when sharing data with others.

Parquet - Binary columnar format, compressed, fast for analytics. Use for large datasets and data pipelines.

Excel - Native spreadsheet format with multiple sheets. Use when stakeholders need to open in Excel directly.

JSON - Structured text format. Use for web APIs or when data has nested structures.

Part 3

Polars: The Modern Alternative

What is Polars?

A **DataFrame library** written in Rust, designed as a faster and more memory-efficient alternative to pandas.

The Problem with Pandas

- **Single-threaded** - Can't use multiple CPU cores
- **High memory usage** - Loads everything into RAM
- **GIL-bound** - Limited by Python's interpreter lock
- **Inconsistent API** - Many ways to do the same thing

How Polars Solves It

- **Multi-threaded** - Uses all cores by default
- **Lazy evaluation** - Builds a plan first, then optimizes before running
- **Rust-powered** - Runs outside Python, bypasses GIL
- **Consistent API** - One clear way to do things

10-100x faster than pandas on large datasets

Lazy evaluation optimizes queries before execution

Handles larger-than-RAM datasets via streaming

Polars vs Pandas: Side by Side

Pandas

```python
import pandas as pd
from pathlib import Path

df = pd.read_csv(Path("data") / "sales.csv")

# Calculate revenue and filter
df["revenue"] = df["quantity"] * df["price"]
result = df[df["revenue"] > 1000]

# Group and aggregate (named aggregation)
summary = df.groupby("product").agg(
    total_revenue=("revenue", "sum"),
    avg_revenue=("revenue", "mean"),
    total_qty=("quantity", "sum")
)
```

Polars

```python
import polars as pl
from pathlib import Path

df = pl.read_csv(Path("data") / "sales.csv")

# Calculate revenue and filter
revenue = pl.col("quantity") * pl.col("price")
result = df.with_columns(
    revenue.alias("revenue")
).filter(pl.col("revenue") > 1000)

# Group and aggregate
summary = df.group_by("product").agg(
    pl.col("revenue").sum().alias("total_revenue"),
    pl.col("revenue").mean().alias("avg_revenue"),
    pl.col("quantity").sum().alias("total_qty")
)
```

Polars syntax is more verbose but explicit - you always know what's happening

Expressions vs Contexts

Expressions describe *what* to do. Unlike pandas (which runs each step immediately), Polars collects expressions and optimizes them before execution.

Expressions (the "what")

```
pl.col("price")              # reference column
pl.col("price") * 1.1        # computation
pl.col("x").sum()            # aggregation
pl.lit("USD")                # literal value
```

Expressions are composable:

```
pl.col("name").str.to_lowercase().str.strip_chars()
pl.col("price").round(2).cast(pl.Int32).alias("cents")
```

Contexts (the "where")

```
df.select( ... )            # choose/transform columns
df.with_columns( ... )      # add new columns
df.filter( ... )            # filter rows
df.group_by( ... ).agg( ... )   # aggregate
```

Contexts accept expressions and define what to do with the result.

Pattern: `df.context(expression)` → `df.select(pl.col("price") * 1.1)`

Expression API in Action

```python
import polars as pl
from pathlib import Path

df = pl.read_csv(Path("data") / "sales_data.csv")

result = df.select(
    pl.col("product"),                          # Select column
    pl.col("price").round(2).alias("rounded_price"),      # Transform + rename
    (pl.col("quantity") * pl.col("price")).alias("revenue"),  # Compute
    pl.col("date").str.to_date().alias("parsed_date"),   # Parse dates
    pl.lit("USD").alias("currency")              # Add constant
)
```

select() - Choose columns

```python
df.select(
    pl.col("name"),
    pl.col("price") * 2
)
```

with_columns() - Add columns

```python
df.with_columns(
    (pl.col("a") + pl.col("b"))
    .alias("sum")
)
```

filter() - Filter rows

```python
df.filter(
    pl.col("price") > 100
)
```

Lazy Evaluation

For large data, Polars can build a query plan first and optimize before running.

Eager (runs immediately)

```
df = pl.read_csv(data)
result = df.filter(pl.col("x") > 0)
```

Each operation executes as it's called. Simple, but loads all data into memory.

Lazy (builds plan first)

```
lf = pl.scan_csv(data)
query = lf.filter(pl.col("x") > 0)
            .select(["a", "b"])
result = query.collect()  # runs here
```

Builds a plan, optimizes it, then executes. Only loads needed columns.

Why Lazy Wins for Large Data

Query Optimization	Memory Efficiency	Smarter Parallelism
Reorders operations, removes unused work	Filters before loading, streams large files	Parallelizes across the entire query plan

Performance Comparison

Task: Calculate revenue and sort by store (3 columns, 1M/10M/100M rows)

Pandas

```
df_pd = pd.DataFrame(data)
df_pd["revenue"] = df_pd["quantity"] * df_pd["price"]
result = df_pd.sort_values(["store_id", "revenue"],
                           ascending=[True, False])
```

Polars

```
df_pl = pl.DataFrame(data)
rev = pl.col("quantity") * pl.col("price")
result = df_pl.with_columns(rev.alias("revenue")).sort(
    ["store_id", "revenue"], descending=[False, True])
```

Rows	Pandas	Polars	Speedup
1 million	0.25s	0.03s	9x
10 million	4.1s	0.38s	11x
100 million	75s	4.8s	16x

The speedup **increases with data size** - Polars shines on large datasets

Common Operations Cheat Sheet

Reading/Writing

```python
from pathlib import Path
data = Path("data")

# Read
df = pl.read_csv(data / "file.csv")
df = pl.read_parquet(data / "file.parquet")
lf = pl.scan_csv(data / "file.csv")  # Lazy

# Write
df.write_csv(data / "out.csv")
df.write_parquet(data / "out.parquet")
```

Filtering

```python
df.filter(pl.col("x") > 10)
df.filter(
    (pl.col("x") > 10) & (pl.col("y") < 5)
)
df.filter(pl.col("name").is_in(["A", "B"]))
```

Aggregations

```python
df.group_by("category").agg(
    pl.col("value").sum().alias("total"),
    pl.col("value").mean().alias("average"),
    pl.len().alias("count"),
    pl.col("id").n_unique().alias("unique_ids")
)
```

Sorting & Limiting

```python
df.sort("price", descending=True)
df.sort(["cat", "price"], descending=[True, False])
df.head(10)
df.tail(5)
```

When to Use Polars vs Pandas

Use Polars When

- Working with **large datasets** (>100MB)
- Performance is critical
- Memory is constrained
- Building **data pipelines**
- Starting a **new project**
- Need **lazy evaluation**

Time to process 111 million rows: 5.2 seconds

Use Pandas When

- Working with **small datasets**
- Need specific **library integrations**
- Team familiarity matters
- Using **legacy code**
- Need **extensive ecosystem** (statsmodels, sklearn)
- Quick exploratory analysis

Tip: Convert between them with Polars: `pl_df.to_pandas()` and `pl.from_pandas(pd_df)`

Part 4

DuckDB: SQL Meets DataFrames

What is DuckDB?

An **in-process SQL database** designed for fast analytical queries. Think "SQLite, but for data analysis."

How It Works

- Runs inside your Python process (no server)
- Columnar storage (reads only needed columns)
- Vectorized execution (operates on entire columns, not row-by-row like traditional databases)
- Queries files directly (CSV, Parquet, JSON) without reading into memory.

Why It Exists

Traditional databases are built for transactions (OLTP). DuckDB is built for **analytics** (OLAP) - aggregations, joins, and scans on large tables.

Key Features

- **SQL interface** - familiar syntax
- **Fast** - optimized for analytics
- **Reads anything** - CSV, Parquet, JSON, pandas, polars
- **Connects anywhere** - PostgreSQL, MySQL, SQLite

```python
import duckdb

result = duckdb.sql("""
    SELECT product, SUM(sales)
    FROM 'sales_data.csv'
    GROUP BY product
""")
```

DuckDB Basics

```python
import duckdb
con = duckdb.connect()                          # In-memory (default)
con = duckdb.connect("data/my_data.duckdb")     # Persistent database
```

Query Files Directly

```python
# CSV
result = con.execute("""
    SELECT * FROM read_csv_auto('data/sales.csv')
    WHERE price > 100
""").df()

# Parquet (supports globs)
result = con.execute("""
    SELECT * FROM 'data/*.parquet'
""").df()

# JSON
result = con.execute("""
    SELECT * FROM read_json_auto('data/records.json')
""").df()
```

Query DataFrames

```python
import pandas as pd
import polars as pl

df_pandas = pd.read_csv("data/sales.csv")
df_polars = pl.read_csv("data/sales.csv")

# Query pandas with SQL
result = con.execute("""
    SELECT product, AVG(price)
    FROM df_pandas
    GROUP BY product
""").df()

# Query polars, return as polars
result = con.execute("""
    SELECT * FROM df_polars WHERE qty > 10
""").pl()
```

SQL for Data Analysis

```python
import duckdb
con = duckdb.connect()

result = con.execute("""
    WITH monthly_sales AS (
        SELECT date_trunc('month', date) as month, product,
               SUM(quantity * price) as revenue, COUNT(*) as transactions
        FROM read_csv_auto('data/sales_data.csv')
        GROUP BY date_trunc('month', date), product
    )
    SELECT *,
           revenue / SUM(revenue) OVER (PARTITION BY month) as pct_of_month
    FROM monthly_sales
    ORDER BY month, revenue DESC
""").df()
```

Standard SQL features work:

- CTEs - `WITH` clauses for multi-step queries
- Window functions - `OVER`, `PARTITION BY`
- Date functions - `date_trunc`, `extract`

Query external databases directly:

```sql
ATTACH 'postgres://user:pass@host/db' AS pg;
SELECT * FROM pg.sales WHERE revenue > 1000;
```

When to Use DuckDB

Perfect For

- **Ad-hoc analysis** on files
- **SQL-first** workflows
- **Large file** processing
- **Joining multiple** data sources
- **Complex aggregations**
- Quick data exploration

Use With Polars When

- Need **programmatic** transformations
- Building **data pipelines**
- **Streaming** processing
- Need **Python ecosystem**
- Complex **business logic**

Avoid When

- Need **OLTP** (transactions)
- High **concurrency** writes
- Need a **server database**
- Real-time **streaming**
- Simple operations on small data

Rule of thumb: If you're working with files larger than your RAM, try DuckDB first

DuckDB + Polars Integration

Best of both worlds: SQL queries with Polars performance

```python
import polars as pl
import duckdb

df = pl.read_parquet("data/large_data.parquet")  # Read with Polars
con = duckdb.connect()

# Use DuckDB for complex SQL on Polars DataFrame
result = con.execute("""
    SELECT store_id, date_trunc('week', timestamp) as week,
           SUM(quantity * price) as weekly_revenue,
           COUNT(DISTINCT customer_id) as unique_customers
    FROM df
    GROUP BY store_id, date_trunc('week', timestamp)
    HAVING SUM(quantity * price) > 10000
""").pl()  # Return as Polars DataFrame

# Continue processing in Polars
final = result.with_columns(
    (pl.col("weekly_revenue") / pl.col("unique_customers")).alias("rev_per_customer")
)
```

Workflow: Load with Polars -> Query with DuckDB SQL -> Process result in Polars

Performance Comparison: Polars vs DuckDB

Task: Calculate revenue and sort by store (3 columns, 1M/10M/100M rows)

Polars

```python
df_pl = pl.DataFrame(data)
rev = pl.col("quantity") * pl.col("price")
result = df_pl.with_columns(
    rev.alias("revenue")
).sort(["store_id", "revenue"],
        descending=[False, True])
```

DuckDB

```python
result = con.execute("""
    SELECT *, quantity * price as revenue
    FROM data
    ORDER BY store_id ASC, revenue DESC
""").pl()
```

Rows	Polars	DuckDB	Notes
1M	0.03s	0.45s	Polars 16x faster
10M	0.37s	0.64s	Polars 1.7x faster
100M	4.9s	9.3s	Polars 1.9x faster

Choose based on: **SQL familiarity** -> DuckDB | **Programmatic pipelines** -> Polars

Summary: Choosing the Right Tool

Tool	Best For	Syntax	Speed
pathlib	File/directory operations	Object-oriented	N/A
Pandas	Small data, ecosystem compatibility	DataFrame API	Baseline
Polars	Large data, pipelines, performance	Expression API	9-16x faster
DuckDB	SQL analysis, file queries, joins	SQL	5-10x faster

Recommended Stack

1. **pathlib** for file operations
2. **Polars** as primary DataFrame library
3. **DuckDB** for complex SQL queries
4. **Pandas** only when needed for compatibility

Quick Decision Guide

- Small data + quick analysis -> **Pandas**
- Large data + pipelines -> **Polars**
- Data larger than RAM + multiple files -> **DuckDB**
- Mix and match as needed!

Hands-on Exercise

Analyze Sales Data with All Three Libraries

Using `data/sales_data.csv`:

1. **Pandas**: Read data, calculate total revenue per product

2. **Polars**: Same analysis - compare the syntax

3. **DuckDB**: Find the best sales day using SQL

```python
# Start here
from pathlib import Path
import pandas as pd
import polars as pl
import duckdb

data_path = Path("data") / "sales_data.csv"

# Your code ...
```

Bonus: Time each approach and compare performance!

Resources

Documentation

- pathlib - docs.python.org
- Pandas - pandas.pydata.org
- Polars - docs.pola.rs
- DuckDB - duckdb.org

Tutorials & Guides

- Polars User Guide
- Pandas to Polars Migration
- Modern Polars (cookbook)

Cheat Sheets

- Pandas Cheat Sheet
- Polars Cheat Sheet (community)

This Tutorial

- Slides - mcda-data-processing.pages.dev
- Code - github.com/smuotoe